

Algorithm Design and Implementation

Principle of Algorithms III

Greedy Algorithms

Guoqiang Li

School of Software, Shanghai Jiao Tong University

Coin Changing

Coin changing

Goal. Given U. S. currency denominations $\{1, 5, 10, 25, 100\}$, devise a method to pay amount to customer using fewest coins.

Example \$34.

Cashier's algorithm. At each iteration, add coin of the largest value that does not take us past the amount to be paid.

Example \$2.89.

Cashier's algorithm

```
Cashiers-Algorithm( $x, c_1, c_2, \dots, c_n$ )
```

Sort n coin denominations so that $0 < c_1 < c_2 < \dots < c_n$;

$S \leftarrow \emptyset$;

while $x > 0$ **do**

$k \leftarrow$ largest coin denomination c_k such that $c_k \leq x$;

if no such k **then** Return no solution;

else

$x \leftarrow x - c_k$;

$S \leftarrow S \cup \{k\}$;

end

end

Return S ;

Quiz

Is the cashier's algorithm optimal?

- A Yes, greedy algorithms are always optimal.
- B Yes, for any set of coin denominations $c_1 < c_2 < \dots < c_n$ provided $c_1 = 1$.
- C Yes, because of special properties of U.S. coin denominations.
- D No.

Cashier's algorithm (for arbitrary coin denominations)

- Q. Is cashier's algorithm optimal for any set of denominations?
- A. No. Consider U.S. postage: 1, 10, 21, 34, 70, 100, 350, 1225, 1500.
- Cashier's algorithm: $\$140 = 100 + 34 + 1 + 1 + 1 + 1 + 1$.
 - Optimal: $\$140 = 70 + 70$.
- A. No. It may not even lead to a feasible solution if $c_1 > 1$. 7, 8, 9.
- Cashier's algorithm: $\$15 = 9 + ?$.
 - Optimal: $\$15 = 7 + 8$.

Properties of any optimal solution (for U.S. coin denominations)



Property. Number of **pennies** ≤ 4 .

Proof. Replace 5 pennies with 1 nickel.

Property. Number of **nickels** ≤ 1 .

Property. Number of **quarters** ≤ 3 .

Properties of any optimal solution (for U.S. coin denominations)



Property. Number of **nickels** + number of **dimes** ≤ 2 .

Proof.

- Recall: ≤ 1 nickel.
- Replace 3 dimes and 0 nickels with 1 quarter and 1 nickel;
- Replace 2 dimes and 1 nickel with 1 quarter.

Optimality of cashier's algorithm (for U.S. coin denominations)

Theorem

Cashier's algorithm is optimal for U.S. coins $\{1, 5, 10, 25, 100\}$.

Proof.

by induction on amount to be paid x

DIY!

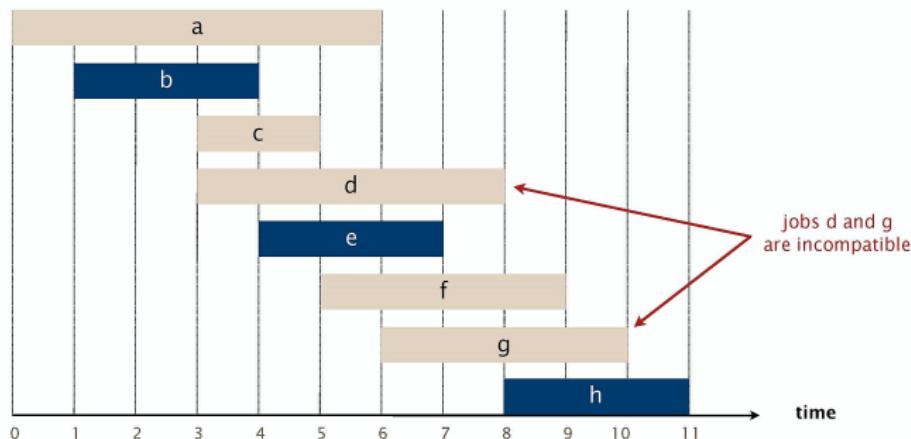
Interval Scheduling

Interval scheduling

Job j starts at s_j and finishes at f_j .

Two jobs compatible if they don't overlap.

Goal: find maximum subset of mutually compatible jobs.



Quiz

Consider jobs in some order, taking each job provided it's compatible with the ones already taken. Which rule is optimal?

- A Earliest start time Consider jobs in ascending order of s_j .
- B Earliest finish time Consider jobs in ascending order of f_j .
- C Shortest interval Consider jobs in ascending order of $f_j - s_j$.
- D None of the above.

Interval scheduling: earliest-finish-time-first algorithm

Earliest-Finish-Time-First($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

Sort jobs by finish times and renumber so that $f_1 \leq f_2 \leq \dots \leq f_n$;

$S \leftarrow \emptyset$;

for $j = 1$ to n **do**

if job j is compatible with S **then**

$| \quad S \leftarrow S \cup \{j\}$

end

end

Return S ;

Proposition. The Earliest-Finish-Time-First algorithm is in $O(n \log n)$ time.

Proof.

- Keep track of job j^* that was added last to S .
- Job j is compatible with S iff $s_j \geq f_{j^*}$.
- Sorting by finish times takes $O(n \log n)$ time.

Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem

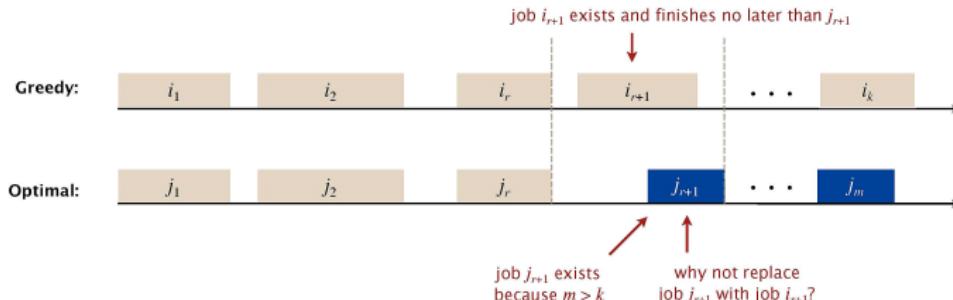
The earliest-finish-time-first algorithm is optimal.

Proof. by contradiction

Assume greedy is not optimal.

Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.

Let j_1, j_2, \dots, j_m denote set of jobs in an optimal solution with $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .



Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem

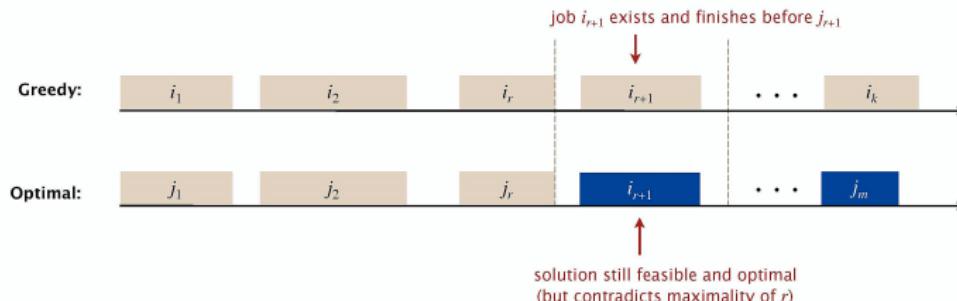
The earliest-finish-time-first algorithm is optimal.

Proof. by contradiction

Assume greedy is not optimal.

Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.

Let j_1, j_2, \dots, j_m denote set of jobs in an optimal solution with $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .



Quiz

Suppose that each job also has a positive weight and the goal is to find a maximum weight subset of mutually compatible intervals. Is the earliest-finish-time-first algorithm still optimal?

- A Yes, because greedy algorithms are always optimal.
- B Yes, because the same proof of correctness is valid.
- C No, because the same proof of correctness is no longer valid.
- D No, because you could assign a huge weight to a job that overlaps the job with the earliest finish time.

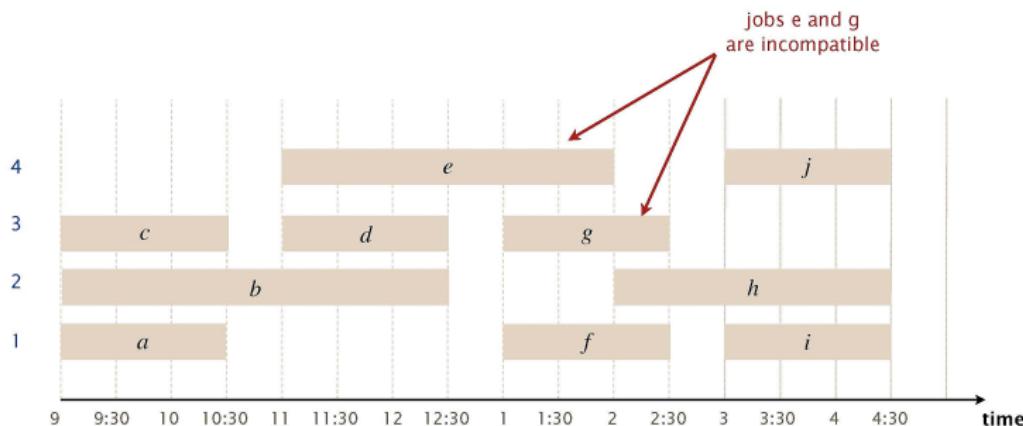
Interval Partitioning

Interval partitioning

Lecture j starts at s_j and finishes at f_j .

Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

Example This schedule uses 4 classrooms to schedule 10 lectures.

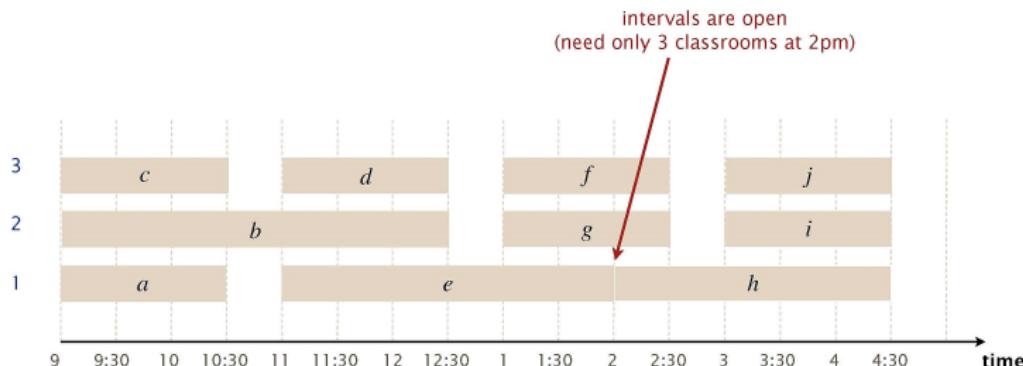


Interval partitioning

Lecture j starts at s_j and finishes at f_j .

Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

Example This schedule uses 3 classrooms to schedule 10 lectures.



Quiz

Consider lectures in some order, assigning each lecture to first available classroom (opening a new classroom if none is available). Which rule is optimal?

- A **Earliest start time** Consider lectures in ascending order of s_j .
- B **Earliest finish time** Consider lectures in ascending order of f_j .
- C **Shortest interval** Consider lectures in ascending order of $f_j - s_j$.
- D None of the above.

Interval partitioning: earliest-start-time-first algorithm

Earliest-Start-Time-First($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

Sort lectures by start times and renumber so that

$s_1 \leq s_2 \leq \dots \leq s_n;$

$d \leftarrow 0;$

for $j = 1$ to n **do**

if lecture j is compatible with some classroom **then**

 | Schedule lecture j in any such classroom k ;

end

else

 | Allocate a new classroom $d + 1$;

 | Schedule lecture j in classroom $d + 1$;

 | $d \leftarrow d + 1$;

end

end

Return schedule;

Interval partitioning: earliest-start-time-first algorithm

Proposition. The earliest-start-time-first algorithm can be implemented in $O(n \log n)$ time.

Proof.

Store classrooms in a **priority queue** (key = finish time of its last lecture).

To determine whether lecture j is compatible with some classroom, compare s_j to key of min classroom k in priority queue.

To add lecture j to classroom k , increase key of classroom k to f_j .

Total number of priority queue operations is $O(n)$.

Sorting by start times takes $O(n \log n)$ time.

Remark. This implementation chooses a classroom k whose finish time of its last lecture is the earliest.

Interval partitioning: lower bound on optimal solution

Definition

The **depth** of a set of open intervals is the maximum number of intervals that contain any given point.

Key observation. Number of classrooms needed \geq depth.

- Q.** Does minimum number of classrooms needed always equal depth?
- A.** Yes! Moreover, earliest-start-time-first algorithm finds a schedule whose number of classrooms equals the depth.

Interval partitioning: analysis of earliest-start-time-first algorithm

Observation. The earliest-start-time first algorithm never schedules two incompatible lectures in the same classroom.

Theorem

Earliest-start-time-first algorithm is optimal.

Proof.

Let d = number of classrooms that the algorithm allocates.

Classroom d is opened because we needed to schedule a lecture, say j , that is incompatible with a lecture in each of $d - 1$ other classrooms.

Thus, these d lectures each end after s_j .

Since we sorted by start time, each of these incompatible lectures start no later than s_j .

Thus, we have d lectures overlapping at time $s_j + \epsilon$.

Scheduling to Minimize Lateness

Scheduling to minimizing lateness

Single resource processes one job at a time.

Job j requires t_j units of processing time and is due at time d_j .

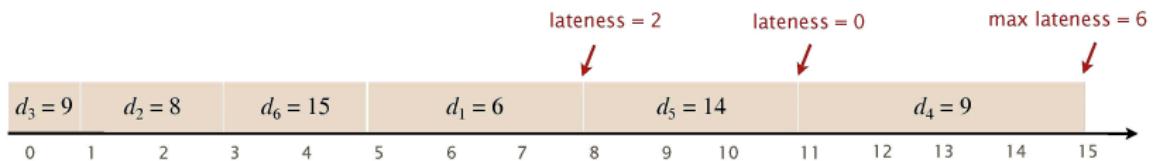
If j starts at time s_j , it finishes at time $f_j = s_j + t_j$.

Lateness: $l_j = \max\{0, f_j - d_j\}$.

Goal: schedule all jobs to minimize maximum lateness $L = \max_j l_j$.

Example

	1	2	3	4	5	6
t_i	3	2	1	4	3	2
d_i	6	8	9	9	14	15



Quiz

Schedule jobs according to some natural order. Which order minimizes the maximum lateness?

- A [shortest processing time] Ascending order of processing time t_j .
- B [earliest deadline first] Ascending order of deadline d_j .
- C [smallest slack] Ascending order of slack: $d_j - t_j$.
- D None of the above.

Minimizing lateness: earliest deadline first

Earliest-Deadline-First($n, t_1, t_2, \dots, t_n, d_1, d_2, \dots, d_n$)

Sort jobs by due times and renumber so that $d_1 \leq d_2 \leq \dots \leq d_n$.

$t \leftarrow 0;$

for $j = 1$ to n **do**

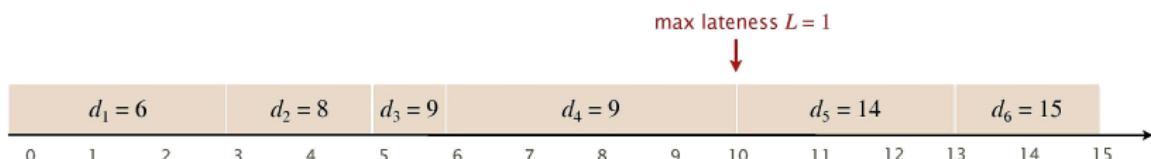
 Assign job j to interval $[t, t + t_j]$;

$s_j \leftarrow t; f_j \leftarrow t + t_j;$

$t \leftarrow t + t_j;$

end

Return intervals $[s_1, f_1], [s_2, f_2], \dots, [s_n, f_n]$;



Minimizing lateness: no idle time

Observation 1. There exists an optimal schedule with no idle time.

Observation 2. The earliest-deadline-first schedule has no idle time.

Minimizing lateness: inversions

Definition (Inversion)

Given a schedule S , an **inversion** is a pair of jobs i and j such that: $i < j$ but j is scheduled before i .

Observation 3. The earliest-deadline-first schedule is the unique idle-free schedule with no inversions.

Minimizing lateness: inversions

Definition

Inversion Given a schedule S , an inversion is a pair of jobs i and j such that: $i < j$ but j is scheduled before i .

Observation 4. If an idle-free schedule has an inversion, then it has an adjacent inversion.

Proof.

Let $i-j$ be a closest inversion.

Let k be element immediately to the right of j .

Case 1 $j > k$ Then $j-k$ is an adjacent inversion.

Case 2 $j < k$ Then $i-k$ is a closer inversion since $i < j < k$.

Minimizing lateness: inversions

Key claim

Exchanging two adjacent, inverted jobs i and j reduces the number of inversions by 1 and does not increase the max lateness.

Proof.

Let l be the lateness before the swap, and let l' be it afterwards.

$$l'_k = l + k \text{ for all } k \neq i, j.$$

$$l'_i \leq l_i$$

If job j is late,

$$\begin{aligned} l'_j &= f'_j - d_j &= f_i - d_j \\ &\leq f_i - d_i &\leq l_i \end{aligned}$$

Minimizing lateness: analysis of earliest-deadline-first algorithm

Theorem

The earliest-deadline-first schedule S is optimal.

Proof. by contradiction

Define S^* to be an optimal schedule with the fewest inversions.

Can assume S^* has no idle time. Observation 1

Case 1. [S^* has no inversions] Then $S = S^*$. Observation 3

Case 2. [S^* has an inversion]

- let $i-j$ be an adjacent inversion Observation 4

- exchanging jobs i and j decreases the number of inversions by 1 without increasing the max lateness Key claim

- contradicts **fewest inversions** part of the definition of S^*

Optimal Caching

Caching.

- Cache with capacity to store k items.
- Sequence of m item requests d_1, d_2, \dots, d_m .
- **Cache hit:** item in cache when requested.
- **Cache miss:** item not in cache when requested.
(must **evict** some item from cache and bring requested item into cache)

Applications. CPU, RAM, hard drive, web, browser, ...

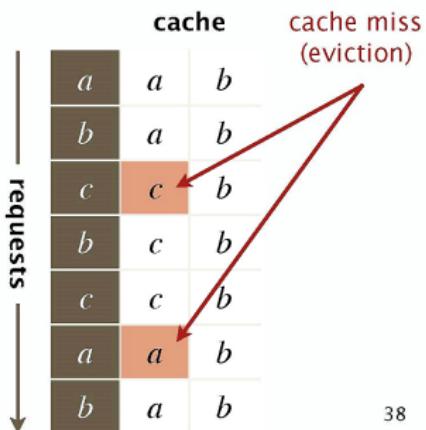
Goal. Eviction schedule that minimizes the number of evictions.

Example

$k = 2$, initial cache = ab , requests: a, b, c, b, c, a, b . Optimal eviction schedule. 2 evictions.

	cache		cache miss (eviction)
	a	b	
a	a	b	
b	a	b	
c	c	b	
b	c	b	
c	c	b	
a	a	b	
b	a	b	

requests →



Optimal offline caching: greedy algorithms

LIFO/FIFO. Evict item brought in least (most) recently.

LRU. Evict item whose most recent access was earliest.

LFU. Evict item that was least frequently requested.

cache						
⋮
a	a	w	x	y	z	FIFO: eject a
d	a	w	x	d	z	LRU: eject d
a	a	w	x	d	z	
b	a	b	x	d	z	
c	a	b	c	d	z	
e	a	b	c	d	e	LIFO: eject e
g	?	?	?	?	?	?
b						
e						
d						
⋮						

A diagram illustrating cache management strategies. A vertical double-headed arrow on the left is labeled "requests" and points downwards. The cache is a grid of 7 columns. Row 1 contains dots. Rows 2 through 8 contain items a through g respectively. Row 9 contains question marks. Row 10 contains b, row 11 e, row 12 d, and row 13 contains dots. Colored cells highlight specific items: a (row 2, column 1), w (row 2, column 2), x (row 2, column 3), y (row 2, column 4), z (row 2, column 5), d (row 4, column 1), b (row 5, column 2), c (row 6, column 3), d (row 6, column 4), e (row 7, column 5), and e (row 8, column 6). Annotations on the right side of the cache explain evictions:

- "FIFO: eject a" is next to the first row of a's.
- "LRU: eject d" is next to the second row of d's.
- "LIFO: eject e" is next to the third row of e's.
- A red arrow points from the text "cache miss (which item to eject?)" to the question marks in row 9.

Optimal offline caching: farthest-in-future (clairvoyant algorithm)

Farthest-in-future. Evict item in the cache that is not requested until farthest in the future.

Theorem (Bélády 1966)

FF is optimal eviction schedule.

Reduced eviction schedules

Definition

A **reduced schedule** is a schedule that brings an item d into the cache in step j only if there is a request for d in step j and d is not already in the cache.

Reduced eviction schedules

Claim

Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.

Proof. by induction on number of steps j

Suppose S brings d into the cache in step j without a request. Let c be the item S evicts when it brings d into the cache.

Case 1a: d evicted before next request for d .

unreduced schedule S			S'		
step j	.	c	.	c	
	.	c	.	c	
	.	c	.	c	
	-d	d	-d	c	
	-d	d	-d	c	
	-d	d	-d	c	
step j'	e	e	e	e	
	e	e	e	e	

d enters cache without a request

d evicted before next request for d

might as well leave c in cache until d is evicted

Reduced eviction schedules

Claim

Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.

Proof. by induction on number of steps j

Suppose S brings d into the cache in step j without a request. Let c be the item S evicts when it brings d into the cache.

Case 1a: d evicted before next request for d .

Case 1b: next request for d occurs before d is evicted.

unreduced schedule S			S'
step j	.	c	
	.	c	
	.	c	
	$\neg d$	d	$\neg d$
	$\neg d$	d	$\neg d$
	$\neg d$	d	$\neg d$
step j'	d	d	d
	.	d	d

d enters cache without a request

d still in cache before next request for d

might as well leave c in cache until d is requested

Reduced eviction schedules

Claim

Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.

Proof. by induction on number of steps j

Suppose S brings d into the cache in step j even though d is in cache. Let c be the item S evicts when it brings d into the cache.

Case 2a: d evicted before it is needed.

unreduced schedule S			S'		
step j	d_1	a	c		
	d_1	a	c		
	d_1	a	c	d_1	a
	d	d_1	a	d₃	
	d	d_1	a	d₃	c
	c	c	a	d₃	
step j'	b	c	a	b	
	d	c	a	d₃	d₃

Annotations:

- d_3 enters cache even though d_1 is already in cache
- d_3 not needed
- d_3 evicted
- d_3 needed
- might as well leave c in cache until d_3 is evicted

Reduced eviction schedules

Claim

Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.

Proof. by induction on number of steps j

Suppose S brings d into the cache in step j even though d is in cache. Let c be the item S evicts when it brings d into the cache.

Case 2a: d evicted before it is needed.

Case 2b: d needed before it is evicted.

unreduced schedule S				S'			
step j	d_1	a	c		d_1	a	c
	d_1	a	c		d_1	a	c
	d_1	a	c	d₃ enters cache even though d₁ is already in cache	d_1	a	c
	d	d_1	a	d₃ not needed	d	d_1	c
	d	d_1	a	d₃ not needed	d	d_1	c
	c	c	a	d_3	c	c	c
	a	c	a	d_3	a	c	c
step j'	d	c	a	d₃ needed	d	c	d₃

Reduced eviction schedules

Claim

Given any unreduced schedule S , can transform it into a reduced schedule S' with no more evictions.

Proof. by induction on number of steps j

- Case 1. S brings d into the cache in step j without a request.
- Case 2. S brings d into the cache in step j even though d is in cache.
- Case 3. If multiple unreduced items in step j , apply each one in turn, dealing with Case 1 before Case 2.

Theorem

FF is optimal eviction algorithm.

Proof. Follows directly from the following invariant.

Farthest-in-future: analysis

Invariant

There exists an optimal reduced schedule S that has the same eviction schedule as S_{FF} through the first j steps.

Proof. by induction on number of steps j

Base case: $j = 0$.

Let S be reduced schedule that satisfies invariant through j steps.

We produce S' that satisfies invariant after $j + 1$ steps.

- Let d denote the item requested in step $j + 1$.
- Since S and S_{FF} have agreed up until now, they have the same cache contents before step $j + 1$.

Invariant

There exists an optimal reduced schedule S that has the same eviction schedule as S_{FF} through the first j steps.

Proof. by induction on number of steps j

Case 1 d is already in the cache.

$S' = S$ satisfies invariant.

Case 2 d is not in the cache and S and S_{FF} evict the same item.

$S' = S$ satisfies invariant.

Invariant

There exists an optimal reduced schedule S that has the same eviction schedule as S_{FF} through the first j steps.

Proof. by induction on number of steps j

Case 3 d is not in the cache; S_{FF} evicts e ; S evicts $f \neq e$.

- begin construction of S' from S by evicting e instead of f
- now S' agrees with S_{FF} for first $j + 1$ steps; we show that having item f in cache is no worse than having item e in cache
- let S' behave the same as S until S' is forced to take a different action (because either S evicts e ; or because either e or f is requested)

Farthest-in-future: analysis

Invariant

There exists an optimal reduced schedule S that has the same eviction schedule as S_{FF} through the first j steps.

Proof. by induction on number of steps j

Let j' be the **first** step after $j+1$ that S' must take a different action from S ; let g denote the item requested in step j' .

Case 3a $g = e$.

Can't happen with FF since there must be a request for f before e .

Case 3b $g \neq f$.

Element f can't be in cache of S ; let e' be the item that S evicts.

- if $e' = e$, S' accesses f from cache; now S and S' have same cache
- if $e' \neq e$, we make S' evict e' and bring e into the cache;
now S and S' have the same cache

We let S' behave exactly like S for remaining requests.

Invariant

There exists an optimal reduced schedule S' that has the same eviction schedule as S_{FF} through the first j steps.

Proof. by induction on number of steps j

Let j' be the **first** step after $j+1$ that S' must take a different action from S ; let g denote the item requested in step j' .

Case 3c $g \neq e, f$. S evicts e .

- make S' evict f .
- now S and S' have the same cache
- let S' behave exactly like S for the remaining requests