

Algorithm Design and Implementation

Principle of Algorithms VII

Divide and Conquer II

Guoqiang Li

School of Software, Shanghai Jiao Tong University

Master Theorem

Divide-and-conquer recurrences

Goal. Recipe for solving common divide-and-conquer recurrences:

$$T(n) = aT\left(\frac{n}{b}\right) + f(n)$$


with $T(0) = 0$ and $T(1) = \Theta(1)$.

Terms.

- $a \geq 1$ is the number of subproblems.
- $b \geq 2$ is the factor by which the subproblem size decreases.
- $f(n) \geq 0$ is the work to divide and combine subproblems.


Recursion tree. [assuming n is a power of b]

- a = branching factor.
- a^i = number of subproblems at level i .
- $1 + \log_b n$ levels.
- n/b^i = size of subproblem at level i .



Divide-and-conquer recurrences: recursion tree

Suppose $T(n)$ satisfies $T(n) = aT(\frac{n}{b}) + n^c$ with $T(1) = 1$, n a power of b .



$$r = a/b^c \quad T(n) = n^c \sum_{i=0}^{\log_b n} r^i$$

Divide-and-conquer recurrences: recursion tree analysis

Suppose $T(n)$ satisfies $T(n) = aT(\frac{n}{b}) + n^c$ with $T(1) = 1$, n a power of b .

Let $r = a/b^c$. Note that $r < 1$ iff $c > \log_b a$.

$$T(n) = n^c \sum_{i=0}^{\log_b n} r^i = \begin{cases} \Theta(n^c) & \text{if } r < 1 \quad c > \log_b a \\ \Theta(n^c \log n) & \text{if } r = 1 \quad c = \log_b a \\ \Theta(n^{\log_b a}) & \text{if } r > 1 \quad c < \log_b a \end{cases}$$

Geometric series.

- If $0 < r < 1$, then $1 + r + r^2 + r^3 + \dots + r^k \leq 1/(1 - r)$.
- If $r = 1$, then $1 + r + r^2 + r^3 + \dots + r^k = k + 1$.
- If $r > 1$, then $1 + r + r^2 + r^3 + \dots + r^k = (r^{k+1} - 1)/(r - 1)$.

Divide-and-conquer recurrences: master theorem

Theorem

Let $a \geq 1$, $b \geq 2$ and $c > 0$ and suppose that $T(n)$ is a function on the non-negative integers that satisfies the recurrence

$$T(n) = aT\left(\frac{n}{b}\right) + \Theta(n^c)$$

with $T(0) = 0$ and $T(1) = \Theta(1)$, where n/b means either $\lfloor n/b \rfloor$ or $\lceil n/b \rceil$. Then,

- If $c < \log_b a$, then $T(n) = \Theta(n^{\log_b a})$.
- If $c = \log_b a$, then $T(n) = \Theta(n^c \log n)$.
- If $c > \log_b a$, then $T(n) = \Theta(n^c)$.

Divide-and-conquer recurrences: master theorem

Theorem

Let $a \geq 1$, $b \geq 2$ and $c > 0$ and suppose that $T(n)$ is a function on the non-negative integers that satisfies the recurrence

$$T(n) = aT\left(\frac{n}{b}\right) + \Theta(n^c)$$

with $T(0) = 0$ and $T(1) = \Theta(1)$, where n/b means either $\lfloor n/b \rfloor$ or $\lceil n/b \rceil$. Then,

- If $c < \log_b a$, then $T(n) = \Theta(n^{\log_b a})$.
- If $c = \log_b a$, then $T(n) = \Theta(n^c \log n)$.
- If $c > \log_b a$, then $T(n) = \Theta(n^c)$.

Extensions.

- Can replace Θ with O everywhere.
- Can replace Θ with Ω everywhere.
- Can replace initial conditions with $T(n) = \Theta(1)$ for all $n \leq n_0$ and require recurrence to hold only for all $n > n_0$.

Divide-and-conquer recurrences: master theorem

Theorem

Let $a \geq 1$, $b \geq 2$ and $c > 0$ and suppose that $T(n)$ is a function on the non-negative integers that satisfies the recurrence

$$T(n) = aT\left(\frac{n}{b}\right) + \Theta(n^c)$$

with $T(0) = 0$ and $T(1) = \Theta(1)$, where n/b means either $\lfloor n/b \rfloor$ or $\lceil n/b \rceil$. Then,

- If $c < \log_b a$, then $T(n) = \Theta(n^{\log_b a})$.
- If $c = \log_b a$, then $T(n) = \Theta(n^c \log n)$.
- If $c > \log_b a$, then $T(n) = \Theta(n^c)$.

Exercise 1. $T(n) = 3T(\lfloor n/2 \rfloor) + 5n$.

- $a = 3, b = 2, c = 1, \log_b a < 1.58$.
- $T(n) = \Theta(n^{\log_2 3}) = O(n^{1.58})$.

Divide-and-conquer recurrences: master theorem

Theorem

Let $a \geq 1$, $b \geq 2$ and $c > 0$ and suppose that $T(n)$ is a function on the non-negative integers that satisfies the recurrence

$$T(n) = aT\left(\frac{n}{b}\right) + \Theta(n^c)$$

with $T(0) = 0$ and $T(1) = \Theta(1)$, where n/b means either $\lfloor n/b \rfloor$ or $\lceil n/b \rceil$. Then,

- If $c < \log_b a$, then $T(n) = \Theta(n^{\log_b a})$.
- If $c = \log_b a$, then $T(n) = \Theta(n^c \log n)$.
- If $c > \log_b a$, then $T(n) = \Theta(n^c)$.

Exercise 2. $T(n) = T(\lfloor n/2 \rfloor) + T(\lceil n/2 \rceil) + 17n$.

- $a = 2, b = 2, c = 1, \log_b a = 1$.
- $T(n) = \Theta(n \log n)$.

Divide-and-conquer recurrences: master theorem

Theorem

Let $a \geq 1$, $b \geq 2$ and $c > 0$ and suppose that $T(n)$ is a function on the non-negative integers that satisfies the recurrence

$$T(n) = aT\left(\frac{n}{b}\right) + \Theta(n^c)$$

with $T(0) = 0$ and $T(1) = \Theta(1)$, where n/b means either $\lfloor n/b \rfloor$ or $\lceil n/b \rceil$. Then,

- If $c < \log_b a$, then $T(n) = \Theta(n^{\log_b a})$.
- If $c = \log_b a$, then $T(n) = \Theta(n^c \log n)$.
- If $c > \log_b a$, then $T(n) = \Theta(n^c)$.

Exercise 3. $T(n) = 48T(\lfloor n/4 \rfloor) + n^3$.

- $a = 48, b = 4, c = 3, \log_b a > 2.79$.
- $T(n) = \Theta(n^3)$.

Quiz 1

Consider the following recurrence. Which case of the master theorem?

$$T(n) = \begin{cases} \Theta(1) & \text{if } n = 1 \\ 3T([n/2]) + \Theta(n) & \text{if } n > 1 \end{cases}$$

- A. Case 1: $T(n) = \Theta(n^{\log_2 3}) = O(n^{1.585})$.
- B. Case 2: $T(n) = \Theta(n \log n)$.
- C. Case 3: $T(n) = \Theta(n)$.
- D. Master theorem not applicable.

Quiz 2

Consider the following recurrence. Which case of the master theorem?

$$T(n) = \begin{cases} 0 & \text{if } n \leq 1 \\ T(\lfloor n/5 \rfloor) + T(n - 3\lfloor n/10 \rfloor) + \frac{11}{5}n & \text{if } n > 1 \end{cases}$$

- A. Case 1: $T(n) = \Theta(n)$.
- B. Case 2: $T(n) = \Theta(n \log n)$.
- C. Case 3: $T(n) = \Theta(n)$.
- D. Master theorem not applicable.

Master theorem need not apply

Gaps in master theorem.

- Number of subproblems is not a constant.

$$T(n) = n T(n/2) + n^2$$

- Number of subproblems is less than 1.

$$T(n) = \frac{1}{2} T(n/2) + n^2$$

- Work to divide and combine subproblems is not $\Theta(n^c)$.

$$T(n) = 2T(n/2) + n \log n$$

Theorem (Akra–Bazzi 1998)

Given constants $a_i > 0$ and $0 < b_i < 1$, functions $|h_i(n)| = O(n/\log^2 n)$ and $g(n) = O(n^c)$. If $T(n)$ satisfies the recurrence:

$$T(n) = \sum_{i=1}^k a_i T(b_i n + h_i(n)) + g(n)$$

then, $T(n) = \Theta\left(n^p \left(1 + \int_1^n \frac{g(u)}{u^{p+1}} du\right)\right)$, where p satisfies $\sum_{i=1}^k a_i b_i^p = 1$.

Example. $T(n) = T(\lfloor n/5 \rfloor) + T(n - 3\lfloor n/10 \rfloor) + 11/5n$, with $T(0) = 0$ and $T(1) = 0$.

- $a_1 = 1, b_1 = 1/5, a_2 = 1, b_2 = 7/10 \Rightarrow p = 0.83978 \dots < 1$.
- $h_1(n) = \lfloor n/5 \rfloor - n/5, h_2(n) = 3/10n - 3\lfloor n/10 \rfloor$.
- $g(n) = 11/5n \Rightarrow T(n) = \Theta(n)$.

Integer Multiplication

Integer addition and subtraction

Addition. Given two n -bit integers a and b , compute $a + b$.

Subtraction. Given two n -bit integers a and b , compute $a - b$.

Grade-school algorithm. $\Theta(n)$ bit operations.

1	1	1	1	1	1	0	1	
	1	1	1	1	1	1	0	1
+	0	1	1	1	1	1	0	1
1	0	1	0	1	0	0	1	0

Remark. Grade-school addition and subtraction algorithms are **optimal**.

Integer multiplication

Multiplication. Given two n -bit integers a and b , compute $a \times b$.

Grade-school algorithm. $\Theta(n^2)$ bit operations.

										1	1	0	1	0	1	0	1
									×	1	1	0	1	0	1	0	1
										1	1	0	1	0	1	0	1
										0	0	0	0	0	0	0	0
								1	1	0	1	0	1	0	1		
						1	1	0	1	0	1	0	1				
				1	1	0	1	0	1	0	1						
			1	1	0	1	0	1	0	1							
		1	1	0	1	0	1	0	1								
	1	1	0	1	0	1	0	1									
	0	0	0	0	0	0	0	1									
0	1	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	1

Conjecture. [Kolmogorov 1956] Grade-school algorithm is optimal.

Theorem. [Karatsuba 1960] Conjecture is false.

Divide-and-conquer multiplication

To multiply two n -bit integers x and y :

- Divide x and y into low- and high-order bits.
- Multiply four $n/2$ -bit integers, recursively.
- Add and shift to obtain result.

$$m = \lceil n/2 \rceil$$

$$a = \lfloor x/2^m \rfloor \quad b = x \bmod 2^m$$

$$c = \lfloor y/2^m \rfloor \quad d = y \bmod 2^m$$

$$xy = (2^m a + b)(2^m c + d) = 2^{2m} ac + 2^m (bc + ad) + bd$$

Example. $x = \underbrace{1000}_a \underbrace{1101}_b \quad y = \underbrace{1110}_c \underbrace{0001}_d$

Multiply(x, y, n)

if $n = 1$ **then** Return $x \times y$;

$m \leftarrow \lceil n/2 \rceil$;

$a \leftarrow \lfloor x/2^m \rfloor$; $b \leftarrow x \bmod 2^m$;

$c \leftarrow \lfloor y/2^m \rfloor$; $d \leftarrow y \bmod 2^m$;

$e \leftarrow \text{Multiply}(a, c, m)$;

$f \leftarrow \text{Multiply}(b, d, m)$;

$g \leftarrow \text{Multiply}(b, c, m)$;

$h \leftarrow \text{Multiply}(a, d, m)$;

Return $2^n e + 2^m(g + h) + f$;

Quiz 3

How many bit operations to multiply two n -bit integers using the divide-and-conquer multiplication algorithm?

$$T(n) = \begin{cases} \Theta(1) & \text{if } n = 1 \\ 4T(\lceil n/2 \rceil) + \Theta(n) & \text{if } n > 1 \end{cases}$$

- A. $T(n) = \Theta(n^{1/2})$
- B. $T(n) = \Theta(n \log n)$.
- C. $T(n) = \Theta(n^{\log_2 3}) = O(n^{1.585})$.
- D. $T(n) = \Theta(n^2)$

To multiply two n -bit integers x and y :

- Divide x and y into low- and high-order bits.
- To compute middle term $bc + ad$, use identity:

$$bc + ad = (a + b)(c + d) - ac - bd$$

- Multiply only three $n/2$ -bit integers, recursively.

Carl Friedrich Gauss(1777-1855) noticed that although the product of two complex numbers

$$(a + bi)(c + di) = ac - bd + (bc + ad)i$$

seems to involve four real-number multiplications, it can in fact be done with just three: ac , bd , and $(a + b)(c + d)$.

To multiply two n -bit integers x and y :

- Divide x and y into low- and high-order bits.
- To compute middle term $bc + ad$, use identity:

$$bc + ad = ac + bd - (a - b)(c - d)$$

- Multiply only three $n/2$ -bit integers, recursively.

Recursive Multiplication

Suppose x and y are two n -integers, and assume for convenience that n is a power of 2.

[Hints: For every n there is an n' with $n \leq n' \leq 2n$, where n' a power of 2.]

$$m = \lceil n/2 \rceil$$

$$a = \lfloor x/2^m \rfloor \quad b = x \bmod 2^m$$

$$c = \lfloor y/2^m \rfloor \quad d = y \bmod 2^m$$

$$x = \underbrace{1000}_a \underbrace{1101}_b$$

$$y = \underbrace{1110}_c \underbrace{0001}_d$$

Gauss's trick

$$\begin{aligned} xy &= (2^m a + b)(2^m c + d) = 2^n ac + 2^m(bc + ad) + bd \\ &= 2^{2m}ac + 2^m((a+b)(c+d) - ac - bd) + bd \end{aligned}$$

Karatsuba's trick

$$\begin{aligned} xy &= (2^m a + b)(2^m c + d) = 2^n ac + 2^m(bc + ad) + bd \\ &= 2^{2m}ac + 2^m(ac + bd - (a-b)(c-d)) + bd \end{aligned}$$

Gauss-Multiply(x, y, n)

if $n = 1$ **then** Return $x \times y$;

$m \leftarrow \lceil n/2 \rceil$;

$a \leftarrow \lfloor x/2^m \rfloor$; $b \leftarrow x \bmod 2^m$;

$c \leftarrow \lfloor y/2^m \rfloor$; $d \leftarrow y \bmod 2^m$;

$e \leftarrow \text{Gauss-Multiply}(a, c, m)$;

$f \leftarrow \text{Gauss-Multiply}(b, d, m)$;

$g \leftarrow \text{Gauss-Multiply}(a + b, c + d, m)$;

Return $2^n e + 2^m(g - e - f) + f$;

Proposition

Gauss(Karatsuba)'s algorithm requires $O(n^{1.585})$ bit operations to multiply two n -bit integers.

Proof. Apply Case 1 of the master theorem to the recurrence:

$$T(n) = \begin{cases} \Theta(1) & \text{if } n = 1 \\ 3T(\lceil n/2 \rceil) + \Theta(n) & \text{if } n > 1 \end{cases}$$

$$\Rightarrow T(n) = \Theta\left(n^{\log_2 3}\right) = O\left(n^{1.585}\right)$$

Practice.

- Use base 32 or 64 (instead of base 2).
- Faster than grade-school algorithm for about 320–640 bits.

Integer arithmetic reductions

Integer multiplication. Given two n -bit integers, compute their product.
bigskip

arithmetic problem	formula	bit complexity
integer multiplication	$a \times b$	$M(n)$
integer square	a^2	$\Theta(M(n))$
integer division	$\lfloor a/b \rfloor, a \bmod b$	$\Theta(M(n))$
integer square root	$\lfloor \sqrt{a} \rfloor$	$\Theta(M(n))$

integer arithmetic problems with the same bit complexity $M(n)$ as integer multiplication

History of asymptotic complexity of integer multiplication

year	algorithm	bit operations
12xx	grade school	$O(n^2)$
1962	Karatsuba–Ofman	$O(n^{1.585})$
1963	Toom-3, Toom-4	$O(n^{1.465}), O(n^{1.404})$
1966	Toom–Cook	$O(n^{1+\epsilon})$
1971	Schönhage–Strassen	$O(n \log n \cdot \log \log n)$
2007	Fürer	$n \log n 2^{O(\log^* n)}$
2018	Harvey–van der Hoeven	$O(n \log n \cdot 2^{2 \lg^* n})$
20XX	???	$O(n)$

number of bit operations to multiply two n -bit integers

Remark. GNU Multiple Precision library uses one of first five algorithms depending on n .

The Fast Fourier Transform

Polynomial multiplication

If $A(x) = a_0 + a_1x + \dots + a_dx^d$ and $B(x) = b_0 + b_1x + \dots + b_dx^d$, their product

$$C(x) = c_0 + c_1x + \dots + c_{2d}x^{2d}$$

has coefficients

$$c_k = a_0b_k + a_1b_{k-1} + \dots + a_kb_0 = \sum_{i=0}^k a_ib_{k-i}$$

where for $i > d$, take a_i and b_i to be zero.

Computing c_k from this formula take $O(k)$ step, and finding all $2d + 1$ coefficients would therefore seem to require $\Theta(d^2)$ time.


Q: Can we do better?

An alternative representation


Fact: A degree- d polynomial is uniquely characterized by its values at any $d + 1$ distinct points.

We can specify a degree- d polynomial $A(x) = a_0 + a_1x + \dots + a_dx^d$ by either of the following:

- Its **coefficients** a_0, a_1, \dots, a_d . (coefficient representation).
- The **values** $A(x_0), A(x_1), \dots, A(x_d)$ (value representation).



An alternative representation



The product $C(x)$ has degree $2d$, it is determined by its value at any $2d + 1$ points.

Its value at any given point z is just $A(z)$ times $B(z)$.

Therefore, polynomial multiplication takes **linear time** in the value representation.

The algorithm

Input: Coefficients of two polynomials, $A(x)$ and $B(x)$, of degree d

Output: Their product $C = A \cdot B$

Selection

Pick some points x_0, x_1, \dots, x_{n-1} , where $n \geq 2d + 1$.

Evaluation

Compute $A(x_0), A(x_1), \dots, A(x_{n-1})$ and $B(x_0), B(x_1), \dots, B(x_{n-1})$.

Multiplication

Compute $C(x_k) = A(x_k)B(x_k)$ for all $k = 0, \dots, n - 1$.

Interpolation

Recover $C(x) = c_0 + c_1x + \dots + c_{2d}x^{2d}$

The **selection** step and the **multiplications** are just linear time:

- In a **typical setting** for polynomial multiplication, the coefficients of the polynomials are real number.
- Moreover, are small enough that basic arithmetic operations take **unit time**.

Evaluating a polynomial of degree $d \leq n$ at a single point takes $O(n)$, and so the baseline for n points is $\Theta(n^2)$.

The **Fast Fourier Transform (FFT)** does it in just $O(n \log n)$ time, for a particularly clever choice of x_0, \dots, x_{n-1} .

Evaluation by divide-and-conquer

Q: How to make it efficient?

First idea, we pick the n points,

$$\pm x_0, \pm x_1, \dots, \pm x_{n/2-1}$$

then the computations required for each $A(x_i)$ and $A(-x_i)$ overlap a lot, because the **even power** of x_i coincide with those of $-x_i$.

We need to split $A(x)$ into its odd and even powers, for instance

$$3 + 4x + 6x^2 + 2x^3 + x^4 + 10x^5 = (3 + 6x^2 + x^4) + x(4 + 2x^2 + 10x^4)$$

More generally

$$A(x) = A_e(x^2) + xA_o(x^2)$$

where $A_e(\cdot)$, with the **even-numbered coefficients**, and $A_o(\cdot)$, with the **odd-numbered coefficients**, are polynomials of degree $\leq n/2 - 1$.

Evaluation by divide-and-conquer

Given paired points $\pm x_i$, the calculations needed for $A(x_i)$ can be recycled toward computing $A(-x_i)$:

$$A(x_i) = A_e(x_i^2) + x_i A_o(x_i^2)$$

$$A(-x_i) = A_e(x_i^2) - x_i A_o(x_i^2)$$

Evaluating $A(x)$ at n paired points $\pm x_0, \dots, \pm x_{n/2-1}$ reduces to evaluating $A_e(x)$ and $A_o(x)$ at just $n/2$ points, $x_0^2, \dots, x_{n/2-1}^2$.

If we could recurse, we would get a divide-and-conquer procedure with running time

$$T(n) = 2T(n/2) + O(n) = O(n \log n)$$

How to choose n points?

Aim: To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ to be themselves **plus-minus pairs**.

Q: How can a square be negative?

- We use **complex numbers**.

At the very bottom of the recursion, we have a **single point**, 1 , in which case the level above it must consist of its square roots, $\pm\sqrt{1} = \pm 1$.

The next level up then has $\pm\sqrt{+1} = \pm 1$, as well as the complex numbers $\pm\sqrt{-1} = \pm i$.

By continuing in this manner, we eventually reach the initial set of n points: the **complex n th** roots of unity, that is the n complex solutions of the equation

$$z^n = 1$$

The n -th complex roots of unity

Solutions to the equation $z^n = 1$

- by the multiplication rules: solutions are $z = (1, \theta)$, for θ a multiple of $2\pi/n$.
- It can be represented as

$$1, \omega, \omega^2, \dots, \omega^{n-1}$$

where

$$\omega = e^{2\pi i/n}$$

For n is even:

- These numbers are plus-minus paired.
- Their squares are the $(n/2)$ -nd roots of unity.

The FFT algorithm

FFT(A, ω)

input : coefficient representation of a polynomial $A(x)$ of degree $\leq n - 1$, where n is a power of 2; ω , an n -th root of unity

output: value representation $A(\omega^0), \dots, A(\omega^{n-1})$

if $\omega = 1$ **then** return $A(1)$;

express $A(x)$ in the form $A_e(x^2) + xA_o(x^2)$;

call FFT(A_e, ω^2) to evaluate A_e at even powers of ω ;

call FFT(A_o, ω^2) to evaluate A_o at even powers of ω ;

for $j = 0$ **to** $n - 1$ **do**

 | compute $A(\omega^j) = A_e(\omega^{2j}) + \omega^j A_o(\omega^{2j})$;

end

return($A(\omega^0), \dots, A(\omega^{n-1})$);

FFT moves from **coefficients** to **values** in time just $O(n \log n)$, when the points $\{x_i\}$ are complex n -th roots of unity $(1, \omega, \omega^2, \dots, \omega^{n-1})$.

That is,

$$\langle value \rangle = \text{FFT}(\langle coefficients \rangle, \omega)$$

We will see that the interpolation can be computed by

$$\langle coefficients \rangle = \frac{1}{n} \text{FFT}(\langle values \rangle, \omega^{-1})$$

A matrix reformation

Let's explicitly set down the relationship between our two representations for a polynomial $A(x)$ of degree $\leq n-1$.

$$\begin{bmatrix} A(x_0) \\ A(x_1) \\ \vdots \\ A(x_{n-1}) \end{bmatrix} = \begin{bmatrix} 1 & x_0 & x_0^2 & \dots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \dots & x_1^{n-1} \\ & & \vdots & & \\ 1 & x_{n-1} & x_{n-1}^2 & \dots & x_{n-1}^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

Let M be the matrix in the middle, which is a **Vandermonde** matrix.

- If x_0, x_1, \dots, x_{n-1} are distinct numbers, then M is invertible.
- **evaluation** is multiplication by M , while **interpolation** is multiplication by M^{-1} .

This **reformulation** of our polynomial operations reveals their essential nature more clearly.

It justifies an assumption that $A(x)$ is **uniquely characterized** by its values at any n points.

Vandermonde matrices also have the distinction of being quicker to invert than more general matrices, in $O(n^2)$ time instead of $O(n^3)$.

However, using this for interpolation would still not be fast enough for us..

Interpolation resolved

In linear algebra terms, the [FFT](#) multiplies an arbitrary n -dimensional vector, which we have been calling the coefficient representation, by the $n \times n$ matrix.

$$M_n(\omega) = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega & \omega^2 & \dots & \omega^{n-1} \\ & & \vdots & & \\ 1 & \omega^j & \omega^{2j} & \dots & \omega^{(n-1)j} \\ & & \vdots & & \\ 1 & \omega^{n-1} & \omega^{2(n-1)} & \dots & \omega^{(n-1)(n-1)} \end{bmatrix}$$

Its (j, k) -th entry (starting row- and column-count at zero) is ω^{jk}

The columns of M are **orthogonal** to each other, which is often called the **Fourier basis**.

The FFT is thus a change of basis, a **rigid rotation**. The inverse of M is the opposite rotation, from the Fourier basis back into the standard basis.

Inversion formula

$$M_n(\omega)^{-1} = \frac{1}{n} M_n(\omega^{-1})$$

Interpolation resolved

Take ω to be $e^{2\pi i/n}$, and think of M as **vectors** in \mathbb{C}^n .

Recall that the **angle** between two vectors $u = (u_0, \dots, u_{n-1})$ and $v = (v_0, \dots, v_{n-1})$ in \mathbb{C}^n is just a **scaling factor** times their **inner product**

$$u \cdot v^* = u_0 v_0^* + u_1 v_1^* + \dots + u_{n-1} v_{n-1}^*$$

where z^* denotes the **complex conjugate** of z .

The above quantity is maximized when the vectors lie in the **same direction** and is zero when the vectors are **orthogonal** to each other.

Interpolation resolved

Lemma

The columns of matrix M are orthogonal to each other.

Proof.

- Take the inner product of any columns j and k of matrix M ,

$$1 + \omega^{j-k} + \omega^{2(j-k)} + \dots + \omega^{(n-1)(j-k)}$$

This is a **geometric series** with first term 1, last term $\omega^{(n-1)(j-k)}$, and ratio ω^{j-k} .

- Therefore, if $j \neq k$, it evaluates to

$$\frac{1 - \omega^{n(j-k)}}{1 - \omega^{j-k}} = 0$$

- If $j = k$, then it evaluates to n .

Corollary

$MM^* = nI$, i.e.,

$$M_n^{-1} = \frac{1}{n} M_n^*$$

The definitive FFT algorithm

The FFT takes as input a vector $a = (a_0, \dots, a_{n-1})$ and a complex number ω whose powers $1, \omega, \omega^2, \dots, \omega^{n-1}$ are the complex n -th roots of unity.

It multiplies vector a by the $n \times n$ matrix $M_n(\omega)$, which has (j, k) -th entry ω^{jk} .

The potential for using divide-and-conquer in this matrix-vector multiplication becomes apparent when M 's columns are segregated into evens and odds.

The product of $M_n(\omega)$ with vector $a = (a_0, \dots, a_{n-1})$, a size- n problem, can be expressed in terms of two size- $n/2$ problems: the product of $M_{n/2}(\omega^2)$ with $(a_0, a_2, \dots, a_{n-2})$ and with $(a_1, a_3, \dots, a_{n-1})$.

This divide-and-conquer strategy leads to the definitive FFT algorithm, whose running time is $T(n) = 2T(n/2) + O(n) = O(n \log n)$.

The general FFT algorithm

FFT(a, ω)

input : An array $a = (a_0, a_1, \dots, a_{n-1})$ for n is a power of 2; ω , an n -th root of unity

output: $M_n(\omega)a$

if $\omega = 1$ **then** return a ;

$(s_0, s_1, \dots, s_{n/2-1}) = \text{FFT}((a_0, a_2, \dots, a_{n-2}), \omega^2)$;

$(s'_0, s'_1, \dots, s'_{n/2-1}) = \text{FFT}((a_1, a_3, \dots, a_{n-1}), \omega^2)$;

for $j = 0$ **to** $n/2 - 1$ **do**

$r_j = s_j + \omega^j s'_j$;
 $r_{j+n/2} = s_j - \omega^j s'_j$;

end

return $(r_0, r_1, \dots, r_{n-1})$;

Top 10 algorithms of the 20th century

1946: The Metropolis Algorithm

1947: Simplex Method

1950: Krylov Subspace Method

1951: The Decompositional Approach to Matrix Computations

1957: The Fortran Optimizing Compiler

1959: QR Algorithm

1962: Quicksort

1965: Fast Fourier Transform

1977: Integer Relation Detection

1987: Fast Multipole Method